

Aim:

To understand the Kubernetes Cluster Architecture, install and Spin Up a Kubernetes Cluster on Linux Machines/Cloud Platforms.

Theory:

Container-based microservices architectures have revolutionized how development and operations teams test and deploy modern software. Containers allow companies to scale and deploy applications more efficiently, but they also introduce new challenges, adding complexity by creating a whole new infrastructure ecosystem.

Today, both large and small software companies are deploying thousands of container instances daily. Managing this level of complexity at scale requires advanced tools. Like Kubernetes.

Originally developed by Google, Kubernetes is an open-source container orchestration platform designed to automate the deployment, scaling, and management of containerized applications. Kubernetes has quickly become the de facto standard for container orchestration and is the flagship project of the Cloud Native Computing Foundation (CNCF), supported by major players like Google, AWS, Microsoft, IBM, Intel, Cisco, and Red Hat.

Kubernetes simplifies the deployment and operation of applications in a microservice architecture by providing an abstraction layer over a group of hosts. This allows development teams to deploy their applications while Kubernetes takes care of key tasks, including:

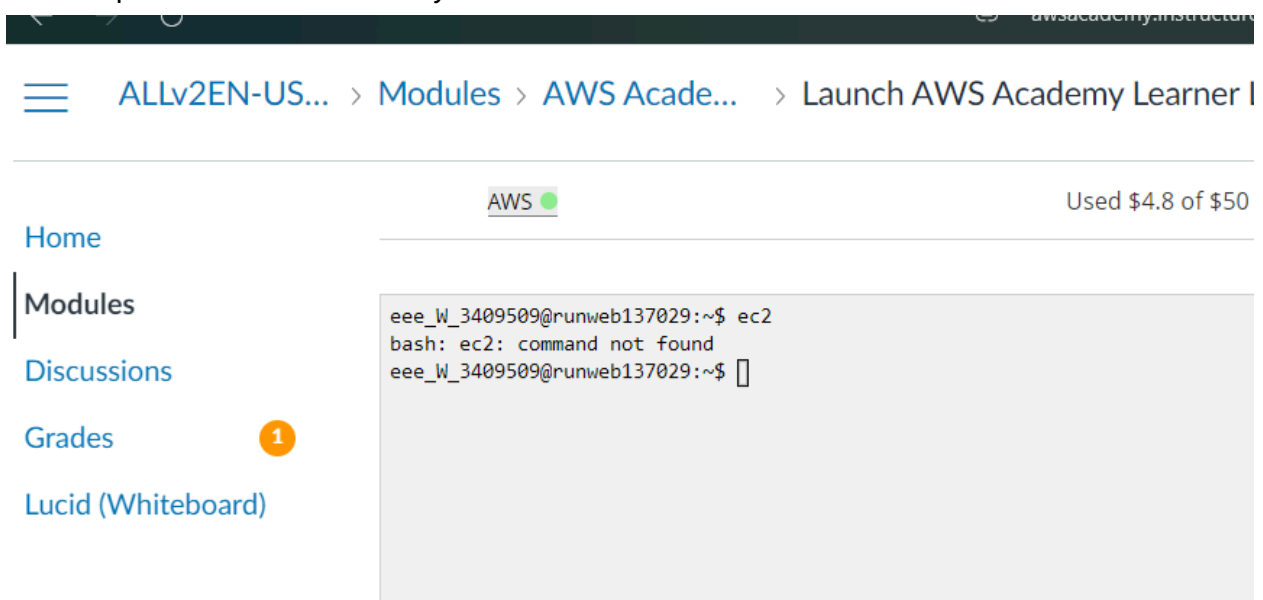
- Managing resource consumption by applications or teams
- Distributing application load evenly across the infrastructure
- Automatically load balancing requests across multiple instances of an application
- Monitoring resource usage to prevent applications from exceeding resource limits and automatically restarting them if needed

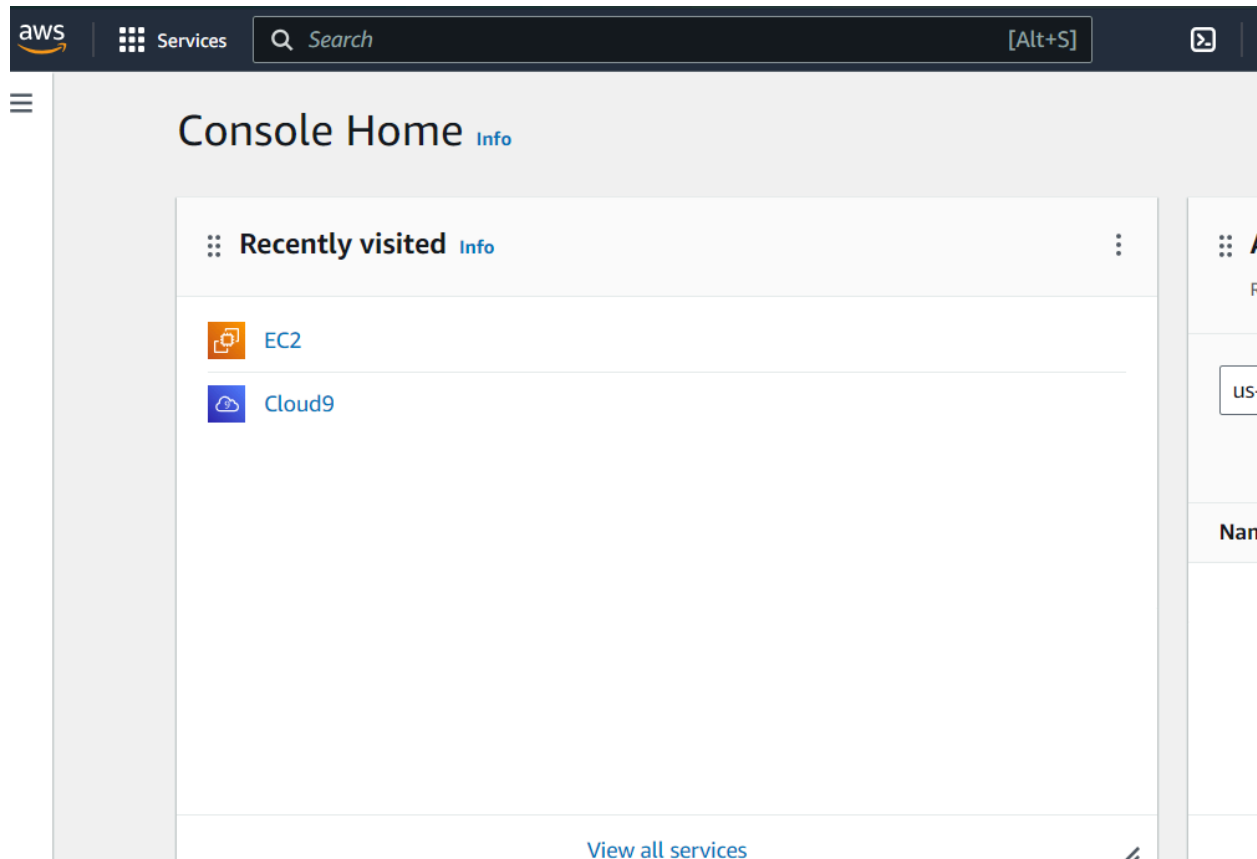
- Moving application instances between hosts when resources are low or if a host fails
- Automatically utilizing additional resources when new hosts are added to the cluster
- Facilitating canary deployments and rollbacks with ease.

Steps:

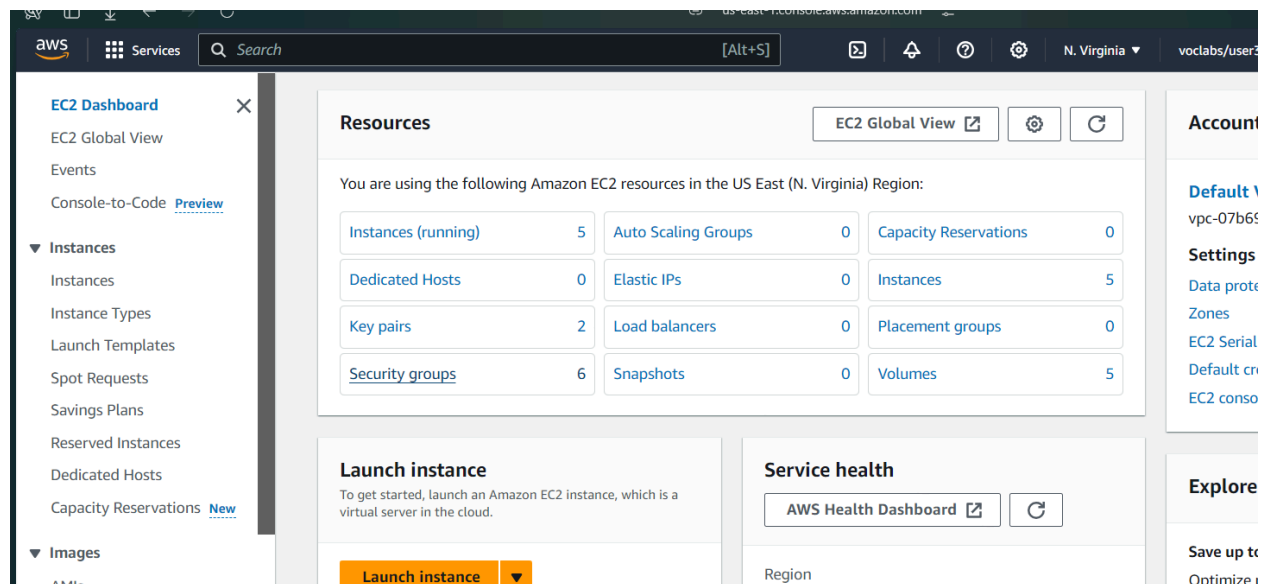
Set Up the instances of each machine

1. open the aws academy.

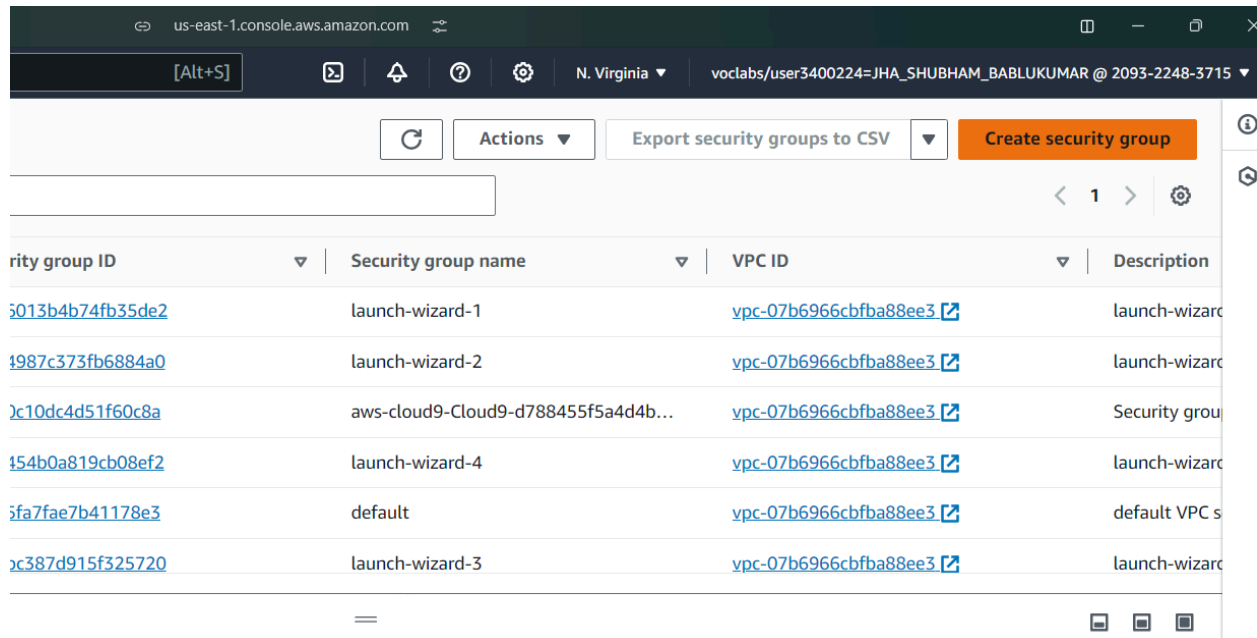




2. Click on security groups



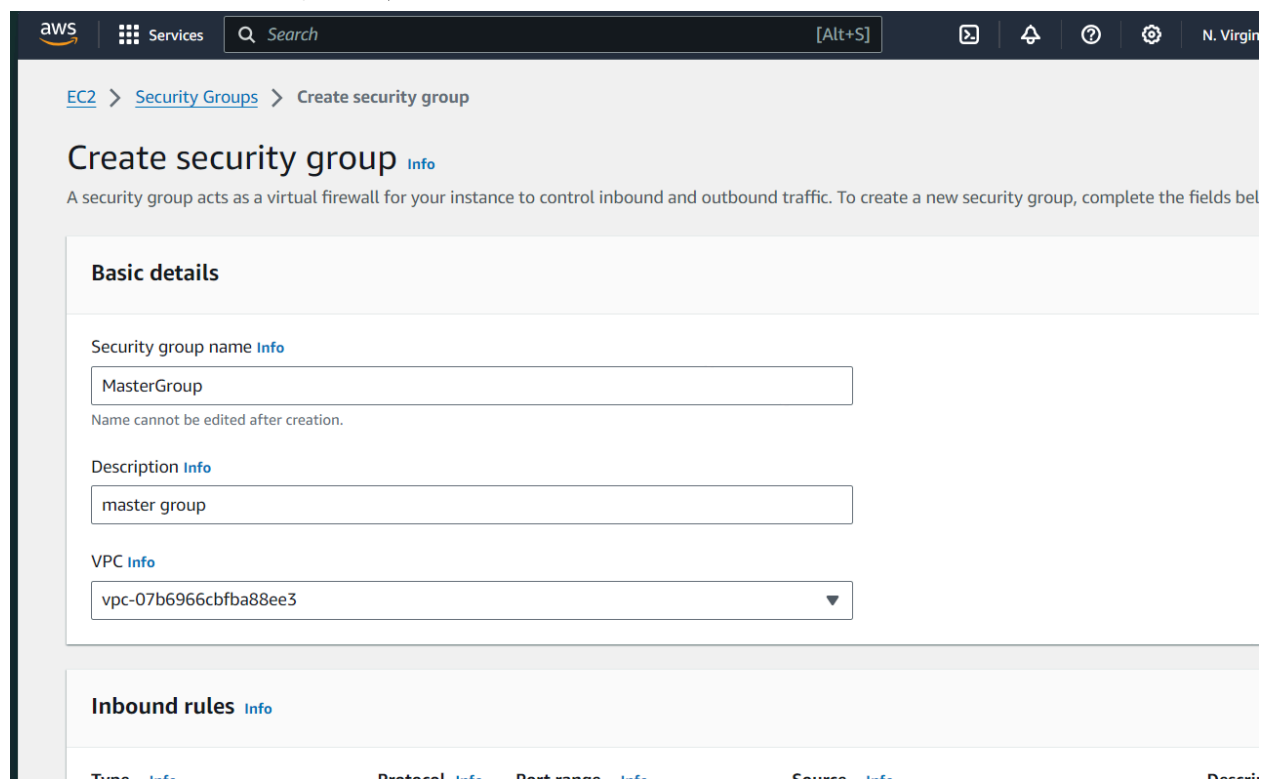
3. create two security groups one for master and other for the two nodes.



The screenshot shows the AWS Management Console for the us-east-1 region. At the top, there are navigation tabs for 'Actions', 'Export security groups to CSV', and a prominent orange 'Create security group' button. Below these is a table listing existing security groups. The table has four columns: 'Security group ID', 'Security group name', 'VPC ID', and 'Description'. There are six rows of data, each representing a different security group, including 'launch-wizard-1' through 'launch-wizard-4', 'aws-cloud9-Cloud9-d788455f5a4d4b...', and 'default'.

Security group ID	Security group name	VPC ID	Description
s013b4b74fb35de2	launch-wizard-1	vpc-07b6966cbfba88ee3	launch-wizard
t987c373fb6884a0	launch-wizard-2	vpc-07b6966cbfba88ee3	launch-wizard
c10dc4d51f60c8a	aws-cloud9-Cloud9-d788455f5a4d4b...	vpc-07b6966cbfba88ee3	Security group
t54b0a819cb08ef2	launch-wizard-4	vpc-07b6966cbfba88ee3	launch-wizard
tfa7fae7b41178e3	default	vpc-07b6966cbfba88ee3	default VPC s
c387d915f325720	launch-wizard-3	vpc-07b6966cbfba88ee3	launch-wizard

4. enter details and add inbound rules (I have given MasterGroup for the master group)



The screenshot shows the 'Create security group' wizard in the AWS Management Console. The page title is 'Create security group' with an 'Info' link. Below the title is a brief description: 'A security group acts as a virtual firewall for your instance to control inbound and outbound traffic. To create a new security group, complete the fields below.' The form is divided into two main sections: 'Basic details' and 'Inbound rules'. The 'Basic details' section contains three fields: 'Security group name' (with the value 'MasterGroup'), 'Description' (with the value 'master group'), and 'VPC' (a dropdown menu showing 'vpc-07b6966cbfba88ee3'). The 'Inbound rules' section is partially visible at the bottom, showing a table with columns for 'Type', 'Protocol', 'Port range', 'Source', and 'Description'.

Create security group [Info](#)

A security group acts as a virtual firewall for your instance to control inbound and outbound traffic. To create a new security group, complete the fields below.

Basic details

Security group name [Info](#)

MasterGroup

Name cannot be edited after creation.

Description [Info](#)

master group

VPC [Info](#)

vpc-07b6966cbfba88ee3

Inbound rules [Info](#)

Type	Protocol	Port range	Source	Description
------	----------	------------	--------	-------------

You have to look for the particular configuration which I did (in the image below)

Inbound rules [Info](#)

Type Info	Protocol Info	Port range Info	Source Info	Description - optional Info	
HTTP	TCP	80	Anywher...	Q 0.0.0.0/0 0.0.0.0/0 X	Delete
Custom TCP	TCP	6443	Anywher...	Q 0.0.0.0/0 0.0.0.0/0 X	Delete
All traffic	All	All	Anywher...	Q 0.0.0.0/0 0.0.0.0/0 X	Delete
Custom TCP	TCP	10251	Anywher...	Q 0.0.0.0/0 0.0.0.0/0 X	Delete
Custom TCP	TCP	10252	Anywher...	Q 0.0.0.0/0 0.0.0.0/0 X	Delete
Custom TCP	TCP	10250	Anywher...	Q 0.0.0.0/0 0.0.0.0/0 X	Delete
All TCP	TCP	0 - 65535	Anywher...	Q 0.0.0.0/0 0.0.0.0/0 X	Delete
SSH	TCP	22	Anywher...	Q 0.0.0.0/0 0.0.0.0/0 X	Delete

click on create security group below.

now do the same for a node group.

Inbound rules [Info](#)

Type Info	Protocol Info	Port range Info	Source Info	Description - optional Info	
All traffic	All	All	Anywher...	Q 0.0.0.0/0 0.0.0.0/0 X	Delete
SSH	TCP	22	Anywher...	Q 0.0.0.0/0 0.0.0.0/0 X	Delete
Custom TCP	TCP	10250	Anywher...	Q 0.0.0.0/0 0.0.0.0/0 X	Delete
All TCP	TCP	0 - 65535	Anywher...	Q 0.0.0.0/0 0.0.0.0/0 X	Delete
Custom TCP	TCP	30000 - 32767	Anywher...	Q 0.0.0.0/0 0.0.0.0/0 X	Delete
HTTP	TCP	80	Anywher...	Q 0.0.0.0/0 0.0.0.0/0 X	Delete

[Add rule](#)

5. now go to ec2 and launch an instance

The screenshot displays the Amazon EC2 console interface. On the left is a navigation sidebar with the following menu items: **EC2 Dashboard** (with a close icon), EC2 Global View, Events, Console-to-Code [Preview](#), **Instances** (expanded), Instances, Instance Types, Launch Templates, Spot Requests, Savings Plans, Reserved Instances, Dedicated Hosts, Capacity Reservations [New](#), **Images** (expanded), AMIs, AMI Catalog, **Elastic Block Store** (expanded), Volumes, Snapshots, and Lifecycle Manager. The main content area is titled 'Resources' and shows a summary of EC2 resources in the US East (N. Virginia) region: 5 running instances, 0 dedicated hosts, 2 key pairs, 8 security groups, 0 auto scaling groups, 0 elastic IPs, 0 load balancers, and 0 snapshots. Below this is the 'Launch instance' section, which includes a description, a prominent orange 'Launch instance' button, a 'Migrate a server' link, and a note about the region. At the bottom, there is an 'Instance alarms' section with a 'View in CloudWatch' link. On the far right, a partial view of the 'Server' section is visible, showing a table with columns for Region, Status, and Zone.

Resources	
Instances (running)	5
Dedicated Hosts	0
Key pairs	2
Security groups	8
Auto Scaling Groups	0
Elastic IPs	0
Load balancers	0
Snapshots	0

Launch instance
To get started, launch an Amazon EC2 instance, which is a virtual server in the cloud.

[Launch instance](#) [Migrate a server](#)

Note: Your instances will launch in the US East (N. Virginia) Region

Instance alarms [View in CloudWatch](#)

add name and set ubuntu:

Amazon EC2 allows you to create virtual machines, or instances, that run on the AWS Cloud. Quickly get started by following the simple steps below.

Name and tags [Info](#)

Name

Master [Add additional tags](#)

▼ Application and OS Images (Amazon Machine Image) [Info](#)

An AMI is a template that contains the software configuration (operating system, application server, and applications) required to launch your instance. Search or Browse for AMIs if you don't see what you are looking for below

Recents | **Quick Start**

Amazon Linux
aws

macOS
Mac

Ubuntu
ubuntu®

Windows
Microsoft

Red Hat
Red Hat

SUSE Linux
SUS

[Browse more AMIs](#)
Including AMIs from AWS, Marketplace and the Community

Amazon Machine Image (AMI)

▼ Summary

Number of instances:

Software Image (AMI): Canonical, Ubuntu 18.04 LTS (64-bit) (ami-0e86e20dae922...)

Virtual server type: t2.micro

Firewall (security group): New security group

Storage (volumes): 1 volume(s) - 8 GiB

Free tier: Includes 750 hours of t2.micro in the Region, 30 million public IPv4 addresses per month, 30 million I/Os, and 100 GB of internet.

create a key if you want

Create key pair

Key pair name

Key pairs allow you to connect to your instance securely.

aws1331

The name can include up to 255 ASCII characters. It can't include leading or trailing spaces.

Key pair type


☒ RSA
RSA encrypted private and public key pair

☐ ED25519
ED25519 encrypted private and public key pair

Private key file format

☒ .pem
For use with OpenSSH

☐ .ppk
For use with PuTTY

 When prompted, store the private key in a secure and accessible location on your computer. **You will need it later to connect to your instance.** [Learn more](#)

Cancel

Create key pair

If you want you can reuse the key pair generated earlier.

Select the security group for master.

No preference (Default subnet in any availability zone)

Auto-assign public IP [Info](#)

Enable

[Additional charges apply](#) when outside of [free tier allowance](#)

Firewall (security groups) [Info](#)

A security group is a set of firewall rules that control the traffic for your instance. Add rules to allow specific traffic to reach your instance.

☐ Create security group ☒ Select existing security group

Common security groups [Info](#)

Select security groups

Q |

<input type="checkbox"/> launch-wizard-1	sg-06013b4b74fb35de2
VPC: vpc-07b6966cbfba88ee3	
<input type="checkbox"/> MasterGroup	sg-00c39d8526dda67f7
VPC: vpc-07b6966cbfba88ee3	
<input type="checkbox"/> launch-wizard-2	sg-04987c373fb6884a0
VPC: vpc-07b6966cbfba88ee3	
aws-cloud9-Cloud9-d788455f5a4d4b4083454091233a80eb- InstanceSecurityGroup-OjiPSymDkJTu	sg-00c10dc4d51f60c8a
VPC: vpc-07b6966cbfba88ee3	
<input type="checkbox"/> launch-wizard-4	sg-0454b0a819cb08ef2
VPC: vpc-07b6966cbfba88ee3	
<input type="checkbox"/> default	sg-05fa7fae7b41178e3
VPC: vpc-07b6966cbfba88ee3	
NodeGroup	sg-0c6df15094da43fbdf

Compare security group rules

erfaces.

Advance

or Magnetic X

with the first 0 instance store

then launch:

The screenshot shows the AWS Management Console interface for launching an EC2 instance. The 'Configure storage' section is active, showing a root volume of 8 GiB with gp3 storage type. A notification indicates that free tier eligible customers can get up to 30 GB of EBS General Purpose (SSD) or Magnetic storage. The 'Advanced details' section is also visible, showing the instance type as t2.micro and the security group as MasterGroup. A 'Launch instance' button is present at the bottom right.

do the same for node instance just select the number of instance as 2.
and select custom security group as node group.

The screenshot shows the AWS Management Console interface for launching an EC2 instance. The 'Network settings' section is active, showing the network as vpc-07b6966cbfa88ee3 and the subnet as No preference. The 'Firewall (security groups)' section shows the 'Select existing security group' option, with 'NodeGroup sg-0cdf15084da43fbdf' selected. The 'Summary' section on the right shows the number of instances as 2, the software image as Canonical, Ubuntu, 24.04, amd64, and the instance type as t2.micro. A 'Launch instance' button is present at the bottom right.

dont give name now. and launch instance.

now go to instances and give name to the blank ones:

Instances (1/8) Info									
Find Instance by attribute or tag (case-sensitive)				All states					
Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4 DNS	Public	
<input type="checkbox"/> Master	i-0ab175e9c60cc3a23	Running	t2.micro	2/2 checks passed	View alarms	us-east-1b	ec2-184-73-147-45.co...	184.7	
<input type="checkbox"/> node-1	i-08ad30b7114767ca2	Running	t2.micro	Initializing	View alarms	us-east-1b	ec2-107-21-179-161.co...	107.2	
<input checked="" type="checkbox"/> node-2	i-03c70d364fb762af5	Running	t2.micro	Initializing	View alarms	us-east-1b	ec2-44-210-122-179.co...	44.21	

6. select master and connect:

Instances (1/8) Info									
Find Instance by attribute or tag (case-sensitive)				All states					
Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4 DNS	Public	
<input checked="" type="checkbox"/> Master	i-0ab175e9c60cc3a23	Running	t2.micro	2/2 checks passed	View alarms	us-east-1b	ec2-184-73-147-45.co...		

click on ssh client:

copy the command below the SSH client session

[EC2](#) > [Instances](#) > [i-0ab175e9c60cc3a23](#) > [Connect to instance](#)

Connect to instance [Info](#)

Connect to your instance i-0ab175e9c60cc3a23 (Master) using any of these options

EC2 Instance Connect | Session Manager | **SSH client** | EC2 serial console

Instance ID

i-0ab175e9c60cc3a23 (Master)

1. Open an SSH client.
2. Locate your private key file. The key used to launch this instance is aws1331.pem
3. Run this command, if necessary, to ensure your key is not publicly viewable.
 `chmod 400 "aws1331.pem"`
4. Connect to your instance using its Public DNS:

`ec2-184-73-147-45.compute-1.amazonaws.com`

Command copied

`ssh -i "aws1331.pem" ubuntu@ec2-184-73-147-45.compute-1.amazonaws.com`

Note: In most cases, the guessed username is correct. However, read your AMI usage instructions to check if the AMI owner has changed the default AMI username.

Cancel

7. Enter the copied command to a terminal window.

```

Microsoft Windows [Version 10.0.22631.4112]
(c) Microsoft Corporation. All rights reserved.

C:\Users\Lenovo>ssh -i "C:\Users\Lenovo\Downloads\aws1331.pem" ubuntu@ec2-34-203-217-53.compute-1.amazonaws.com
The authenticity of host 'ec2-34-203-217-53.compute-1.amazonaws.com (34.203.217.53)' can't be established.
ED25519 key fingerprint is SHA256:3onu4BDyF+uS+Fwt16U1L99+0SyVYZbTNPitWoY074Y.
This key is not known by any other names
Are you sure you want to continue connecting (yes/no/[fingerprint])? |

```

It would prompt whether we want to continue connecting. type yes.

```

System information as of Fri Sep 27 15:25:58 UTC 2024

System load:  0.0          Processes:      104
Usage of /:   22.8% of 6.71GB Users logged in:  0
Memory usage: 19%         IPv4 address for enX0: 172.31.87.211
Swap usage:   0%

Expanded Security Maintenance for Applications is not enabled.

Updates can be applied immediately.

Enable ESM Apps to receive additional future security updates.
See https://ubuntu.com/esm or run: sudo pro status

The list of available updates is more than a week old.
To check for new updates run: sudo apt update

The programs included with the Ubuntu system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*/copyright.

Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by
applicable law.

To run a command as administrator (user "root"), use "sudo <command>".
See "man sudo_root" for details.

ubuntu@ip-172-31-87-211:~$ |

```

The step is similar for node 1 and node 2 instances too. Just use different terminal windows.

node 1:

EC2 > Instances > i-08ad30b7114767ca2 > Connect to instance

Connect to instance [Info](#)

Connect to your instance i-08ad30b7114767ca2 (node-1) using any of these options

EC2 Instance Connect

Session Manager

SSH client

EC2 serial console

Instance ID

i-08ad30b7114767ca2 (node-1)

1. Open an SSH client.
2. Locate your private key file. The key used to launch this instance is aws1331.pem
3. Run this command, if necessary, to ensure your key is not publicly viewable.
 `chmod 400 "aws1331.pem"`
4. Connect to your instance using its Public DNS:
 `ec2-34-88-249-77.compute-1.amazonaws.com`

Command copied

```

System load:  0.0          Processes:      104
Usage of /:   22.8% of 6.71GB Users logged in:  0
Memory usage: 19%         IPv4 address for enX0: 172.31.89.24
Swap usage:   0%

Expanded Security Maintenance for Applications is not enabled.

Updates can be applied immediately.

Enable ESM Apps to receive additional future security updates.
See https://ubuntu.com/esm or run: sudo pro status

The list of available updates is more than a week old.
To check for new updates run: sudo apt update

The programs included with the Ubuntu system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*/copyright.

Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by

```



node 2:

Connect to your instance i-03c70d364fb762af5 (node-2) using any of these options


EC2 Instance Connect Session Manager **SSH client** EC2 serial console


Instance ID

 i-03c70d364fb762af5 (node-2)

1. Open an SSH client.
2. Locate your private key file. The key used to launch this instance is aws1331.pem
3. Run this command, if necessary, to ensure your key is not publicly viewable.
 `chmod 400 "aws1331.pem"`
4. Connect to your instance using its Public DNS:
 `ec2-35-173-124-11.compute-1.amazonaws.com`

 Command copied

 `ssh -i "aws1331.pem" ubuntu@ec2-35-173-124-11.compute-1.amazonaws.com`

 **Note:** In most cases, the guessed username is correct. However, read your AMI usage instructions to check if the AMI owner has changed the default AMI username.

```
System load:  0.08      Processes:    104
Usage of /:   22.8% of 6.71GB   Users logged in:  0
Memory usage: 19%          IPv4 address for enx0: 172.31.88.60
Swap usage:   0%
```

Expanded Security Maintenance for Applications is not enabled.

0 updates can be applied immediately.

Enable ESM Apps to receive additional future security updates.
See <https://ubuntu.com/esm> or run: `sudo pro status`

The list of available updates is more than a week old.
To check for new updates run: `sudo apt update`

The programs included with the Ubuntu system are free software;
the exact distribution terms for each program are described in the
individual files in `/usr/share/doc/*/copyright`.

Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by
applicable law.

To run a command as administrator (user "root"), use "`sudo <command>`".
See "`man sudo_root`" for details.

```
ubuntu@ip-172-31-88-60:~$ |
```

8. From now on run the commands on all the 3 terminals unless instructed otherwise.

and the images (screen shots) will only be of master unless stated otherwise.

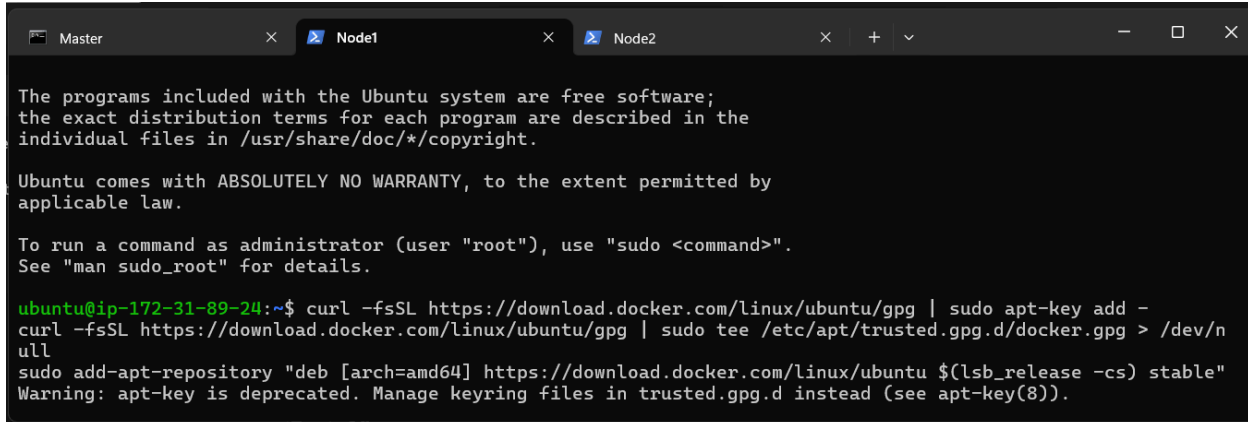
```
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo apt-key add -
```

```
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo tee /etc/apt/trusted.gpg.d/docker.gpg > /dev/null
```

```
sudo add-apt-repository "deb [arch=amd64] https://download.docker.com/linux/ubuntu $(lsb_release -cs) stable"
```

```
ubuntu@ip-172-31-87-211:~$ curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo apt-key add -
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo tee /etc/apt/trusted.gpg.d/docker.gpg > /dev/null
sudo add-apt-repository "deb [arch=amd64] https://download.docker.com/linux/ubuntu $(lsb_release -cs) stable"
Warning: apt-key is deprecated. Manage keyring files in trusted.gpg.d instead (see apt-key(8)).
OK
Repository: 'deb [arch=amd64] https://download.docker.com/linux/ubuntu noble stable'
Description:
Archive for codename: noble components: stable
More info: https://download.docker.com/linux/ubuntu
Adding repository.
Press [ENTER] to continue or Ctrl-c to cancel.
Adding deb entry to /etc/apt/sources.list.d/archive_uri-https_download_docker_com_linux_ubuntu-noble.list
Adding disabled deb-src entry to /etc/apt/sources.list.d/archive_uri-https_download_docker_com_linux_ubuntu-noble.list
Hit:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble InRelease
Get:2 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-updates InRelease [126 kB]
Get:3 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-backports InRelease [126 kB]
Get:4 http://security.ubuntu.com/ubuntu noble-security InRelease [126 kB]
```

now for node1 and node2:



```
Master Node1 Node2
The programs included with the Ubuntu system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*/copyright.

Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by
applicable law.

To run a command as administrator (user "root"), use "sudo <command>".
See "man sudo_root" for details.

ubuntu@ip-172-31-89-24:~$ curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo apt-key add -
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo tee /etc/apt/trusted.gpg.d/docker.gpg > /dev/n
ull
sudo add-apt-repository "deb [arch=amd64] https://download.docker.com/linux/ubuntu $(lsb_release -cs) stable"
Warning: apt-key is deprecated. Manage keyring files in trusted.gpg.d instead (see apt-key(8)).
```

```
Master x Node1 x Node2 x + - □ x
applicable law.

To run a command as administrator (user "root"), use "sudo <command>".
See "man sudo_root" for details.

ubuntu@ip-172-31-88-60:~$ curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo apt-key add -
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo tee /etc/apt/trusted.gpg.d/docker.gpg > /dev/n
ull
sudo add-apt-repository "deb [arch=amd64] https://download.docker.com/linux/ubuntu $(lsb_release -cs) stable"

Warning: apt-key is deprecated. Manage keyring files in trusted.gpg.d instead (see apt-key(8)).
OK
Repository: 'deb [arch=amd64] https://download.docker.com/linux/ubuntu noble stable'
Description:
Archive for codename: noble components: stable
More info: https://download.docker.com/linux/ubuntu
```

sudo apt-get update

sudo apt-get install -y docker-ce

```
Processing triggers for man-db (2.12.0-4build2) ...
Processing triggers for libc-bin (2.39-0ubuntu8.2) ...
Scanning processes...
Scanning linux images...

Running kernel seems to be up-to-date.

No services need to be restarted.

No containers need to be restarted.

No user sessions are running outdated binaries.

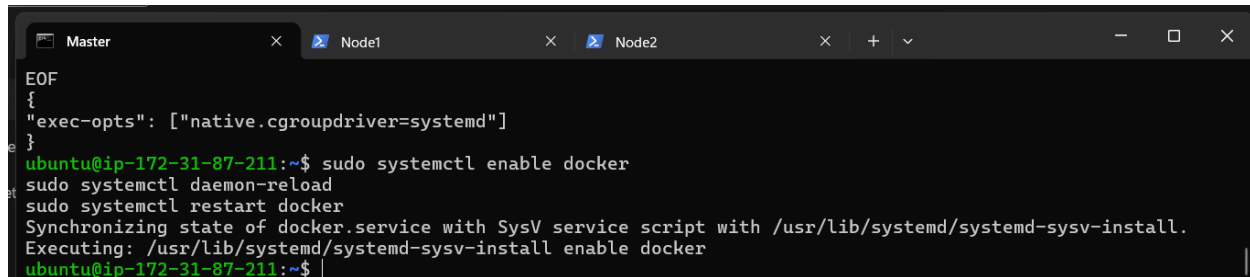
No VM guests are running outdated hypervisor (qemu) binaries on this host.
ubuntu@ip-172-31-87-211:~$ |
```

```
No VM guests are running outdated hypervisor (qemu) binaries on this host.
ubuntu@ip-172-31-87-211:~$ sudo mkdir -p /etc/docker
cat <<EOF | sudo tee /etc/docker/daemon.json
{
  "exec-opts": ["native.cgroupdriver=systemd"]
}
EOF
{
  "exec-opts": ["native.cgroupdriver=systemd"]
}
ubuntu@ip-172-31-87-211:~$ |
```

sudo systemctl enable docker

```
sudo systemctl daemon-reload
```

```
sudo systemctl restart docker
```

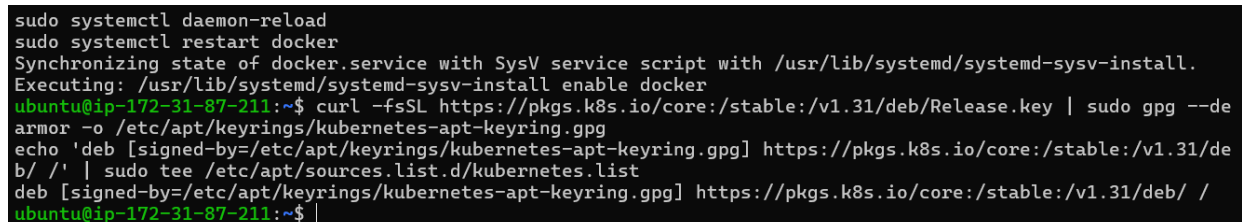
A terminal window with three tabs: 'Master', 'Node1', and 'Node2'. The 'Node1' tab is active. The terminal shows the following commands and output:

```
EOF
{
"exec-opts": ["native.cgroupdriver=systemd"]
}
ubuntu@ip-172-31-87-211:~$ sudo systemctl enable docker
sudo systemctl daemon-reload
sudo systemctl restart docker
Synchronizing state of docker.service with SysV service script with /usr/lib/systemd/systemd-sysv-install.
Executing: /usr/lib/systemd/systemd-sysv-install enable docker
ubuntu@ip-172-31-87-211:~$
```

9. Run the below command to install Kubernetes.

```
curl -fsSL https://pkgs.k8s.io/core:/stable:/v1.31/deb/Release.key |
sudo gpg --dearmor -o /etc/apt/keyrings/kubernetes-apt-keyring.gpg
```

```
echo 'deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg]
https://pkgs.k8s.io/core:/stable:/v1.31/deb/ /' | sudo tee
/etc/apt/sources.list.d/kubernetes.list
```

A terminal window showing the continuation of the previous commands:

```
sudo systemctl daemon-reload
sudo systemctl restart docker
Synchronizing state of docker.service with SysV service script with /usr/lib/systemd/systemd-sysv-install.
Executing: /usr/lib/systemd/systemd-sysv-install enable docker
ubuntu@ip-172-31-87-211:~$ curl -fsSL https://pkgs.k8s.io/core:/stable:/v1.31/deb/Release.key | sudo gpg --de
armor -o /etc/apt/keyrings/kubernetes-apt-keyring.gpg
echo 'deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg] https://pkgs.k8s.io/core:/stable:/v1.31/de
b/ /' | sudo tee /etc/apt/sources.list.d/kubernetes.list
deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg] https://pkgs.k8s.io/core:/stable:/v1.31/deb/ /
ubuntu@ip-172-31-87-211:~$
```

Run the commands:

```
sudo apt-get update
```

```
sudo apt-get install -y kubelet kubeadm kubectl
```

```
sudo apt-mark hold kubelet kubeadm kubectl
```



```
ubuntu@ip-172-31-87-211:~$ sudo apt-get update
sudo apt-get install -y kubelet kubeadm kubectl
sudo apt-mark hold kubelet kubeadm kubectl
Hit:1 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble InRelease
Get:2 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-updates InRelease [126 kB]
Hit:3 http://us-east-1.ec2.archive.ubuntu.com/ubuntu noble-backports InRelease
Hit:4 http://security.ubuntu.com/ubuntu noble-security InRelease
Hit:5 https://download.docker.com/linux/ubuntu noble InRelease
Get:6 https://prod-cdn.packages.k8s.io/repositories/isv:/kubernetes:/core:/stable:/v1.31/deb InRelease [1186 B]
Get:7 https://prod-cdn.packages.k8s.io/repositories/isv:/kubernetes:/core:/stable:/v1.31/deb Packages [4865 B]
```

```
No services need to be restarted.

No containers need to be restarted.

No user sessions are running outdated binaries.

No VM guests are running outdated hypervisor (qemu) binaries on this host.
kubelet set on hold.
kubeadm set on hold.
kubectl set on hold.
```

sudo systemctl enable --now kubelet

sudo apt-get install -y containerd

```
ubuntu@ip-172-31-87-211:~$ sudo systemctl enable --now kubelet
sudo apt-get install -y containerd
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following packages were automatically installed and are no longer required:
  docker-buildx-plugin docker-ce-cli docker-ce-rootless-extras docker-compose-plugin libltdl7
  pigz slirp4netns
Use 'sudo apt autoremove' to remove them.
```

```
Running kernel seems to be up-to-date.

No services need to be restarted.

No containers need to be restarted.

No user sessions are running outdated binaries.

No VM guests are running outdated hypervisor (qemu) binaries on this host.
ubuntu@ip-172-31-87-211:~$ |
```

sudo mkdir -p /etc/containerd

sudo containerd config default | sudo tee /etc/containerd/config.toml

```
Master Node1 Node2
No services need to be restarted.
No containers need to be restarted.
No user sessions are running outdated binaries.
No VM guests are running outdated hypervisor (qemu) binaries on this host.
ubuntu@ip-172-31-87-211:~$ sudo mkdir -p /etc/containerd
sudo containerd config default | sudo tee /etc/containerd/config.toml
disabled_plugins = []
imports = []
oom_score = 0
plugin_dir = ""
required_plugins = []
root = "/var/lib/containerd"
state = "/run/containerd"
temp = ""

[timeouts]
"io.containerd.timeout.bolt.open" = "0s"
"io.containerd.timeout.metrics.shimstats" = "2s"
"io.containerd.timeout.shim.cleanup" = "5s"
"io.containerd.timeout.shim.load" = "5s"
"io.containerd.timeout.shim.shutdown" = "3s"
"io.containerd.timeout.task.state" = "2s"

[ttrpc]
address = ""
gid = 0
uid = 0
ubuntu@ip-172-31-87-211:~$ |
```

`sudo systemctl restart containerd`

`sudo systemctl enable containerd`

`sudo systemctl status containerd`

```
Sep 27 16:24:28 ip-172-31-87-211 containerd[4900]: time="2024-09-27T16:24:28.133071347Z" level=info msg=serve>
Sep 27 16:24:28 ip-172-31-87-211 containerd[4900]: time="2024-09-27T16:24:28.133103323Z" level=info msg=serve>
Sep 27 16:24:28 ip-172-31-87-211 containerd[4900]: time="2024-09-27T16:24:28.133174859Z" level=info msg="Sta>
Sep 27 16:24:28 ip-172-31-87-211 containerd[4900]: time="2024-09-27T16:24:28.133199320Z" level=info msg="Sta>
Sep 27 16:24:28 ip-172-31-87-211 containerd[4900]: time="2024-09-27T16:24:28.133236780Z" level=info msg="Sta>
Sep 27 16:24:28 ip-172-31-87-211 containerd[4900]: time="2024-09-27T16:24:28.133244763Z" level=info msg="Sta>
Sep 27 16:24:28 ip-172-31-87-211 containerd[4900]: time="2024-09-27T16:24:28.133252776Z" level=info msg="Sta>
ubuntu@ip-172-31-87-211:~$ |
```

exit with `ctrl+c`.

sudo apt-get install -y socat

```
ubuntu@ip-172-31-87-211:~$ sudo apt-get install -y socat
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following packages were automatically installed and are no longer required:
  docker-buildx-plugin docker-ce-cli docker-ce-rootless-extras docker-compose-plugin libltdl7
  pigz slurp4netns
Use 'sudo apt autoremove' to remove them.
```

Run the following command in master only:

sudo kubeadm init --pod-network-cidr=10.244.0.0/16

if it gives error use:

**sudo kubeadm init --pod-network-cidr=10.244.0.0/16
--ignore-preflight-errors=NumCPU,Mem**

```
To see the stack trace of this error execute with --v=5 or higher
ubuntu@ip-172-31-87-211:~$ sudo kubeadm init --pod-network-cidr=10.244.0.0/16 --ignore-preflight-errors=NumCPU,Mem
[init] Using Kubernetes version: v1.31.0
[preflight] Running pre-flight checks
[WARNING NumCPU]: the number of available CPUs 1 is less than the required 2
[WARNING Mem]: the system RAM (957 MB) is less than the minimum 1700 MB
[preflight] Pulling images required for setting up a Kubernetes cluster
[preflight] This might take a minute or two, depending on the speed of your internet connection
[preflight] You can also perform this action beforehand using 'kubeadm config images pull'
W0927 16:47:12.068193 6025 checks.go:846] detected that the sandbox image "registry.k8s.io/pause:3.8" of the container runtime is inconsistent with that used by kubeadm. It is recommended to use "registry.k8s.io/pause:3.10" as the CRI sandbox image.
```

Alternatively, if you are the root user, you can run:

```
export KUBECONFIG=/etc/kubernetes/admin.conf
```

You should now deploy a pod network to the cluster.

Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
<https://kubernetes.io/docs/concepts/cluster-administration/addons/>

Then you can join any number of worker nodes by running the following on each as root:

```
kubeadm join 172.31.87.211:6443 --token lhbvbx.kzlxu87vmv8hvx4o \
  --discovery-token-ca-cert-hash sha256:1cef7709c45a42691a2ff0e44e3acf7f0e214fec7f4f822bb6818f3cfd24ea4
3
ubuntu@ip-172-31-87-211:~$
```

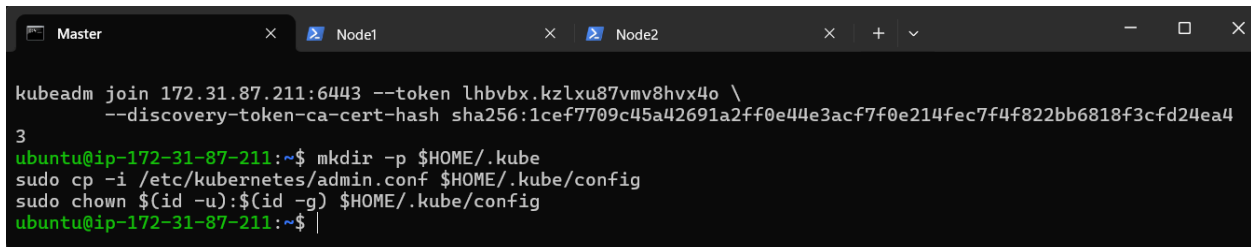
Token and ca

Note: copy the text after kubeadm that you see at the later part like below:

```
kubeadm join 172.31.87.211:6443 --token lhbvbx.kzlXu87vmv8hvx4o \  
--discovery-token-ca-cert-hash sha256:1cef7709c45a42691a2ff0e44e3acf7f0e214fec7f4f822bb6818f3cfd24ea43
```

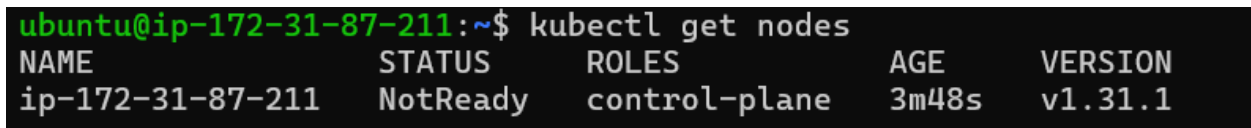
Run this command on master

```
mkdir -p $HOME/.kube  
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config  
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```



```
Master Node1 Node2  
kubeadm join 172.31.87.211:6443 --token lhbvbx.kzlXu87vmv8hvx4o \  
--discovery-token-ca-cert-hash sha256:1cef7709c45a42691a2ff0e44e3acf7f0e214fec7f4f822bb6818f3cfd24ea4  
3  
ubuntu@ip-172-31-87-211:~$ mkdir -p $HOME/.kube  
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config  
sudo chown $(id -u):$(id -g) $HOME/.kube/config  
ubuntu@ip-172-31-87-211:~$
```

Now Run the command `kubectl get nodes` to see the nodes before executing Join command on nodes.



```
ubuntu@ip-172-31-87-211:~$ kubectl get nodes  
NAME STATUS ROLES AGE VERSION  
ip-172-31-87-211 NotReady control-plane 3m48s v1.31.1
```

Now paste the token and ca that I asked to copy earlier, on both the nodes.

use sudo before them.

it would be something like:

```
sudo kubeadm join <your-master-node-ip>:6443 --token <your-token>
--discovery-token-ca-cert-hash sha256:<your-ca-cert-hash>
(it has placeholders)
```

Node1:

```
ubuntu@ip-172-31-89-24:~$ sudo kubeadm join 172.31.87.211:6443 --token lhbvbx.kzlxu87vmv8hvx4o \
--discovery-token-ca-cert-hash sha256:1cef7709c45a42691a2ff0e44e3acf7f0e214fec7f4f822bb6818f3cfd24ea4
3
[preflight] Running pre-flight checks
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -o yaml'
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-check] Waiting for a healthy kubelet at http://127.0.0.1:10248/healthz. This can take up to 4m0s
```

Node2:

```
ubuntu@ip-172-31-88-60:~$ sudo kubeadm join 172.31.87.211:6443 --token lhbvbx.kzlxu87vmv8hvx4o \
--discovery-token-ca-cert-hash sha256:1cef7709c45a42691a2ff0e44e3acf7f0e214fec7f4f822bb6818f3cfd24ea4
3
[preflight] Running pre-flight checks
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -o yaml'
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-check] Waiting for a healthy kubelet at http://127.0.0.1:10248/healthz. This can take up to 4m0s
[kubelet-check] The kubelet is healthy after 1.00236733s
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap
```

Step 9: Now Run the command on Master **kubectl get nodes** to see the nodes after executing Join command on nodes.

```
ubuntu@ip-172-31-87-211:~$ kubectl get nodes
NAME                STATUS    ROLES    AGE   VERSION
ip-172-31-87-211    NotReady control-plane 7m35s v1.31.1
ip-172-31-88-60     NotReady <none>    12s   v1.31.1
ip-172-31-89-24     NotReady <none>    32s   v1.31.1
ubuntu@ip-172-31-87-211:~$ |
```

Step 10: Since Status is NotReady we have to add a network plugin. And also we have to give

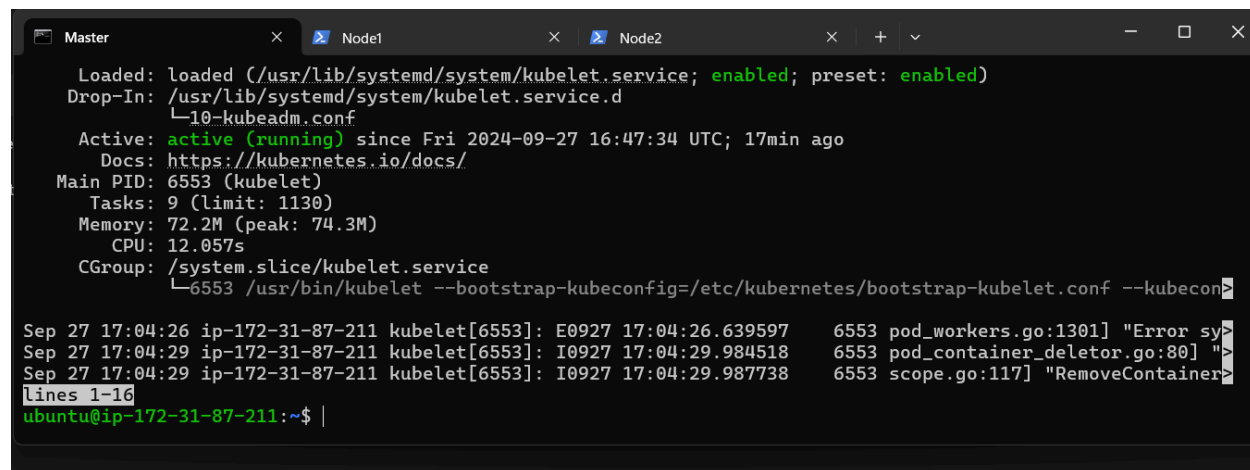
the name to the nodes.

paste this command on master terminal. (and the following commands will be based on master unless states otherwise.

kubectl apply -f <https://docs.projectcalico.org/manifests/calico.yaml>

sudo systemctl status kubelet

again use ctrl+c to exit.



```
Master x Node1 x Node2 + v - □ x
Loaded: loaded (/usr/lib/systemd/system/kubelet.service; enabled; preset: enabled)
Drop-In: /usr/lib/systemd/system/kubelet.service.d
└─10-kubeadm.conf
Active: active (running) since Fri 2024-09-27 16:47:34 UTC; 17min ago
Docs: https://kubernetes.io/docs/
Main PID: 6553 (kubelet)
Tasks: 9 (limit: 1130)
Memory: 72.2M (peak: 74.3M)
CPU: 12.057s
CGroup: /system.slice/kubelet.service
└─6553 /usr/bin/kubelet --bootstrap-kubeconfig=/etc/kubernetes/bootstrap-kubelet.conf --kubecon>
Sep 27 17:04:26 ip-172-31-87-211 kubelet[6553]: E0927 17:04:26.639597 6553 pod_workers.go:1301] "Error sy>
Sep 27 17:04:29 ip-172-31-87-211 kubelet[6553]: I0927 17:04:29.984518 6553 pod_container_deletor.go:80] ">
Sep 27 17:04:29 ip-172-31-87-211 kubelet[6553]: I0927 17:04:29.987738 6553 scope.go:117] "RemoveContainer>
lines 1-16
ubuntu@ip-172-31-87-211:~$ |
```

Now Run command **kubectl get nodes -o wide** we can see Status is ready.

```

lines 1-16
ubuntu@ip-172-31-87-211:~$ kubectl get nodes -o wide
NAME                                STATUS    ROLES    AGE    VERSION    INTERNAL-IP    EXTERNAL-IP    OS-IMAGE
Kubernetes-VERSION    CONTAINER-RUNTIME
ip-172-31-87-211      Ready    control-plane    18m    v1.31.1    172.31.87.211    <none>        Ubuntu 24.04 LTS
6.8.0-1012-aws        containerd://1.7.12
ip-172-31-88-60      Ready    <none>    11m    v1.31.1    172.31.88.60    <none>        Ubuntu 24.04 LTS
6.8.0-1012-aws        containerd://1.7.12
ip-172-31-89-24      Ready    <none>    11m    v1.31.1    172.31.89.24    <none>        Ubuntu 24.04 LTS
6.8.0-1012-aws        containerd://1.7.12
ubuntu@ip-172-31-87-211:~$ |

```

Now to Rename run this command

Syntax: `kubectl label node <node-ip> kubernetes.io/role=worker`

examples:

Rename to Node 1: `kubectl label node ip-<node1ip>
kubernetes.io/role=Node1`

Rename to Node 2: `kubectl label node ip-<node2ip>
kubernetes.io/role=Node2`

```

ubuntu@ip-172-31-87-211:~$ kubectl label node ip-172-31-88-60 kubernetes.io/role=Node2
node/ip-172-31-88-60 labeled
ubuntu@ip-172-31-87-211:~$ kubectl label node ip-172-31-89-24 kubernetes.io/role=Node1
node/ip-172-31-89-24 labeled
ubuntu@ip-172-31-87-211:~$ |

```

Step 11: Run command `kubectl get nodes -o wide` . And Hence we can see we have Successfully connected Node 1 and Node 2 to the Master.

```

node/ip-172-31-89-24 labeled
ubuntu@ip-172-31-87-211:~$ kubectl get nodes -o wide
NAME                                STATUS    ROLES    AGE    VERSION    INTERNAL-IP    EXTERNAL-IP    OS-IMAGE
Kubernetes-VERSION    CONTAINER-RUNTIME
ip-172-31-87-211      Ready    control-plane    24m    v1.31.1    172.31.87.211    <none>        Ubuntu 24.04 LTS
6.8.0-1012-aws        containerd://1.7.12
ip-172-31-88-60      Ready    Node2    17m    v1.31.1    172.31.88.60    <none>        Ubuntu 24.04 LTS
6.8.0-1012-aws        containerd://1.7.12
ip-172-31-89-24      Ready    Node1    17m    v1.31.1    172.31.89.24    <none>        Ubuntu 24.04 LTS
6.8.0-1012-aws        containerd://1.7.12
ubuntu@ip-172-31-87-211:~$ |

```

Conclusion:

In this Advanced DevOps Lab experiment, we began by setting up three EC2 Ubuntu instances

on AWS, designating one as the Master node and the others as Worker nodes.

We then installed Docker and Kubernetes on all instances, ensuring Docker was properly configured.

The Kubernetes cluster was initialized on the Master node, and the Flannel networking plugin

was applied to facilitate communication between nodes.

Finally, we joined the Worker nodes to the cluster using the provided token and hash, resulting in a fully operational Kubernetes cluster ready for managing and scaling containerized applications.